CONSTRUCT DIGITAL · INTERNAL EXPLAINER

How Kunal runs BD and accounts with AI

Ten specialist agents, one orchestrator, and a strict set of rules — turning Claude from a capable chatbot into an operational layer that produces Construct-quality output consistently.

April 2026 • For Charan + team • Written by Kunal

USE CASE 1

"Hey Con" — the morning check-in

I open my laptop, type **"hey con"**, and go make coffee. By the time I'm back, the whole picture of my day is on screen — what's urgent, what's changed since yesterday, what's blocked, what I missed. No typing. No re-briefing.

K

Kunal types "hey con"

Two words. That's the whole prompt.

0

Sync with remote

Con pulls latest from GitHub — catches any commits from other sessions or machines.

1

Priority list + live calendar

Reads the live priority list. Queries Google Calendar for today + next 2 days.

— Fires Friday → Littlefinger revenue pulse
— Fires Monday → Fury weekly performance review
— Fires quarter-start → memory audit

2

Outside-in sweep (the truth check)

Con trusts live data more than stored notes. Runs in parallel:

— Gmail inbox — unread, important, addressed to me
— Gmail sent — what I've already actioned in the last few days
— Slack — every active channel, my @mentions, messages I sent

3

Task log reconciliation

Cross-checks the task log against what the sweep actually found. Live data always wins — task log updated on the spot.

— Overdue items flagged (promised date passed, no movement)
— Stale client contexts flagged (>14 days since touch)
— Open assumptions closed if resolved

4

Inbox / outbox triage — zero practice

Every file in either folder gets exactly one recommended action: CATALOGUE ARCHIVE MERGE DELETE MOVE KUNAL ACTION
Nothing sits without a decision attached.

5

Con reports — verified state only

Briefing on screen: today's RED items from priority list + calendar, what changed since last session, blocked/overdue, pipeline pulse, Kunal's open action items from outbox.

6

Ask Kunal

Up to 5 focused questions for things only I have the answer to — WhatsApp conversations, phone calls, things outside Gmail/Slack. Short, specific.

From "hey con" to a complete verified briefing — usually under 90 seconds. The same protocol runs every single session. No drift, no skipped checks.

USE CASE 2

New BD inquiry → first-touch email draft

A new inbound inquiry lands in my email — a prospective client via web form, referral, or LinkedIn. I paste it into chat. Within about a minute, I have a researched brief plus a first-touch email written in my voice, ready to review and send. **Six agents touch it before it reaches me.**

K

Kunal pastes the inquiry

The raw inbound — subject, sender, CC line, body. No summary, no framing, just paste.

C

Con receives + declares dispatch

Reads the inquiry, identifies the domain, declares: *"DISPATCHING TO: Varys → Hormozi"*.

M

MYCROFT **Pre-task context check**

Does a client folder already exist? Has this company contacted us before? Any prior context? If yes, loads it. If no, notes clean slate. Packages context for the specialists.

V

VARYS **Research brief**

Fetches the company website, checks the contact on LinkedIn, scans for recent news.

— Company profile — what they do, sector, audience
— Contact — seniority, function, likely decision weight
— Signals — budget hints, process stage, CC line
— Relevant Construct proof points — matched case studies
— Flags — wrong-company risks, unreadable websites

Each section rated High / Medium / Unverified confidence.

H

HORMOZI **BD brief + email draft**

Consumes Varys's brief. Reads the email drafting guidelines. Creates `clients/[name]/context.md`. Produces:

— Structured new-business brief (our POV, open questions, suggested engagement)
— First-touch email draft in Kunal's voice — no signature, ready to review

Never answers their RFP point-by-point. Signals fit, proposes call, asks for materials.

C

Con reviews

Does the email sound like Kunal, not a BD coordinator? Does it match inquiry depth (detailed vs blank)? Are proof points accurate? If no → re-brief Hormozi. If yes → packages for Kunal.

K

Kunal approves / edits / sends

Reviews brief + email. Usually 1–2 small tweaks. Sends from my Gmail. Total time from paste to send: ~5 minutes.

M

MYCROFT **Post-task capture**

Updates `clients/[name]/context.md` with what was sent. Logs to GDrive engagement log. Updates memory pointers. Nothing to remember later — it's all captured.

C

Con logs performance

Rates Varys and Hormozi on this task (1–5). Fury reads these weekly to sharpen the profiles over time.

Six agents, one orchestrator, one scripted protocol. Every inbound inquiry runs through the same shape — which is why the output is consistent.

THE ROSTER

Ten specialists, each with a persona + a role

Persona names are mnemonics — the role is what matters. Each agent has a detailed profile with skills, tools, quality standards, and the specific Construct files they read before every task.

Con

ORCHESTRATOR

Routes every task, briefs specialists, reviews outputs. Never executes himself.

Alex Hormozi

NEW BUSINESS

Handles inbound BD inquiries — brief, research synthesis, first-touch email draft in Kunal's voice.

Alfred

ACCOUNTS

Manages active clients from signed SOW onward. Relationship emails, context maintenance, account health.

Hemingway

PROPOSALS & SOWS

Writes commercial documents. Clean prose, no placeholders, no boilerplate errors. Authors slide content plans for Da Vinci to build.

Da Vinci

DESIGN & BRAND

Builds decks via pptxgenjs. Enforces brand compliance — logo sizing, colour rotation, block layouts.

Varys

RESEARCH

Prospect, contact, and competitive intelligence. Mandatory on every BD outbound and proposal.

Nick Fury

RECRUITMENT

Hires team members and designs new agents. Maintains the roster and performance reviews.

Mycroft

CONTEXT & MEMORY (PROTOCOL LAYER)

Runs automatically before and after every task. Loads prior context in, captures new context out.

Chanakya

STRATEGY

Q

AUTOMATION

Commercial judgement — proposal angle, pricing floor, pursuit/withdraw calls, portfolio prioritisation.

Builds and maintains n8n workflows. BD pipeline infrastructure, inbox notifications, event triggers.

Littlefinger

REVENUE

ARF, pipeline coverage, cold-deal detection, resource forecast. Delivers the Friday revenue pulse proactively.

Dewey*

HOUSEKEEPING & FILE REGISTRY

Inbox/outbox zero enforcement, orphan file detection, weekly drift audits, file registry maintenance. (*candidate)

WHY THIS BEATS PLAIN CLAUDE

The delta — same AI, different structure

1 **Consistent Construct-specific output**

Hemingway reads `sow/SKILL.md` before writing — our template, our payment structure, our case study library, our banned phrases. First draft lands 90% there, not 50%.

2 **Role separation — no competence blur**

Hemingway writes. Da Vinci designs. One agent can't be excellent at both commercial writing AND pptxgenjs AND brand compliance. Splitting them preserves quality on both sides.

3 **Hard dispatch rule as a brake**

"Con never self-executes" forces decomposition. Extra 10 seconds of "which agent owns this?" usually saves 30 minutes of revision.

4 **Institutional memory across sessions**

Mycroft reads prior client context automatically. I never re-brief from scratch. System picks up where it left off.

5 **Clean handoffs between stages**

Inquiry → Varys → Hormozi → Hemingway → Da Vinci → (signed SOW) → Alfred. Each agent knows what the previous one produced. Nothing falls through the gap.

6 **Performance compounds**

Every completed task gets rated. Fury reviews weekly. Profiles get sharpened on real data. The system gets better over time.

HONEST READ

What this doesn't do

- **Input quality still matters.** Vague brief, vague output.
- **WhatsApp + phone calls are invisible.** Any conversation outside Gmail/Slack needs me to surface it.
- **Agents don't replace the team.** Charan, Luis, Zulayha still do the actual creative, dev, paid media work. This layer helps me run BD + account coordination faster.
- **New territory is slow at first.** The first time we do a new client type or service line, output is mediocre. By the fifth, the pattern is captured in a SKILL file and subsequent output is consistent.
- **Rules need maintenance.** Agent files and SKILL files need to be kept current. Quarterly audit cadence is in place to catch drift.

The bottom line: same underlying AI. Dramatically more useful because of the structure around it.

Full details and all agent profiles in the Construct repo → `con/agents/ROSTER.md`

Happy to walk through any part of this live if useful.

— Kunal, April 2026